

# Deep Generative Models

## Lecture 1

Roman Isachenko



Ozon Masters

Spring, 2021

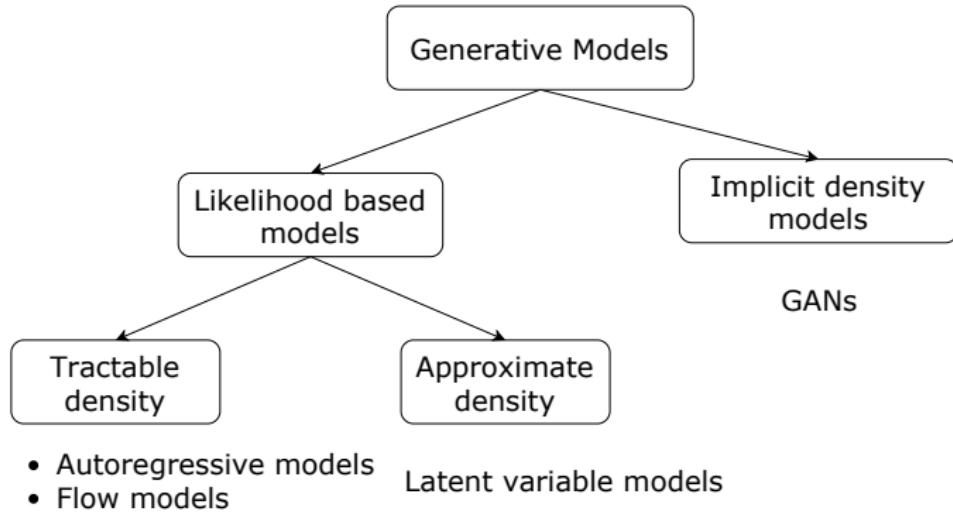
# Logistics

- ▶ homeworks: 30 points
  - ▶ hw1: autoregressive models
  - ▶ hw2: latent variable models
  - ▶ hw3: flow models
  - ▶ hw4: adversarial models
- ▶ exam: 30 points
- ▶ final project: 40 points

Last year course page: [link](#)

Admission: [link](#)

# Generative models zoo

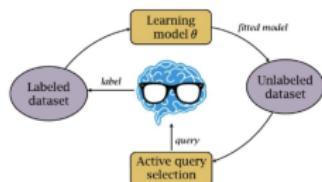


# Applications

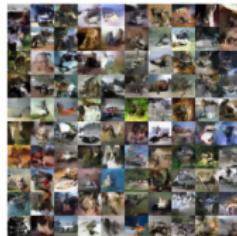
" i want to talk to you . "  
" i want to be with you . "  
" i do n't want to be with you . "  
" i do n't want to be with you . "  
she did n't want to be with him .

he was silent for a long moment .  
he was silent for a moment .  
it was quiet for a moment .  
it was dark and cold .  
there was a pause .  
it was my turn .

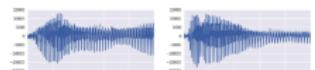
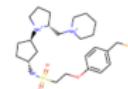
## Text analysis



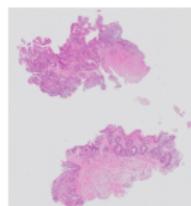
## Active Learning



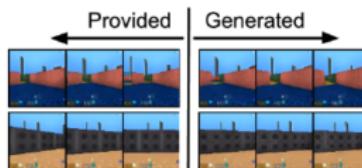
## Image analysis



## Audio analysis



## Medical data



## Reinforcement Learning

and more...

## Applications: Image generation (VAE)



# Applications: Image generation (DCGAN)



---

Radford A., Metz L., Chintala S. *Unsupervised representation learning with deep convolutional generative adversarial networks*, 2015

# Applications: SuperResolution (SRGAN)



---

Ledig C. et al. Photo-realistic single image super-resolution using a generative adversarial network, 2016

# Applications: Face generation (StyleGAN)



Karras T., Laine S., Aila T. A style-based generator architecture for generative adversarial networks, 2018

## Applications: Face generation (VQ-VAE-2)



---

Razavi A., Oord A., Vinyals O. Generating Diverse High-Fidelity Images with VQ-VAE-2, 2019

# Applications: Language Modelling

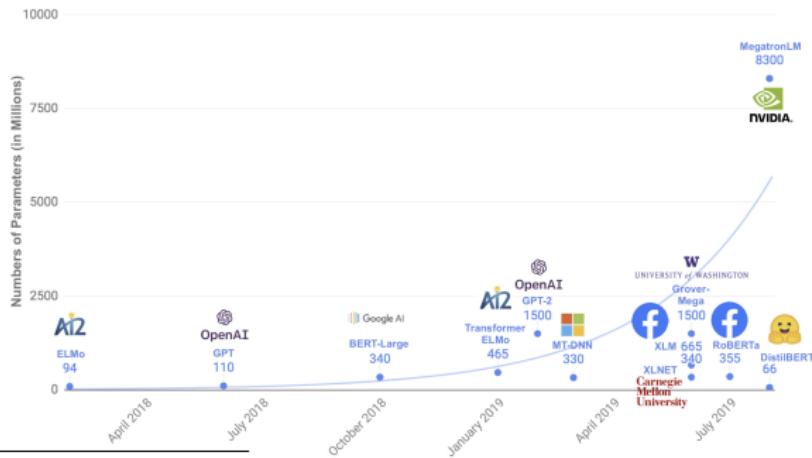


image credit: <http://jalammar.github.io/illustrated-gpt2>

*Sanh V. et al. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter, 2019.*

# Problem Statement

We are given i.i.d. samples  $\{\mathbf{x}_i\}_{i=1}^n \in X$  (e.g.  $X = \mathbb{R}^m$ ) from unknown distribution  $\pi(\mathbf{x})$ .

## Goal

We would like to learn a distribution  $\pi(\mathbf{x})$  for

- ▶ evaluating  $\pi(\mathbf{x})$  for new samples (how likely to get object  $\mathbf{x}$ ?);
- ▶ sampling from  $\pi(\mathbf{x})$  (to get new objects  $\mathbf{x} \sim \pi(\mathbf{x})$ ).

## Challenge

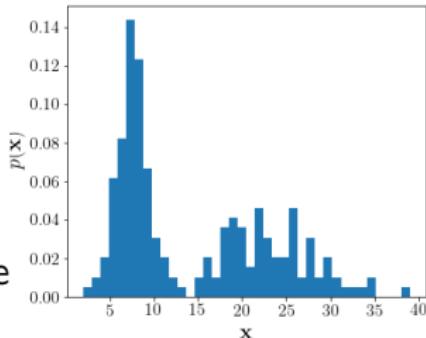
Data is complex and high-dimensional. Imagine the dataset of images which lies in the space  $X \subset \mathbb{R}^{\text{width} \times \text{height}}$ .

## Histogram as a generative model

Let  $x \sim \text{Categorical}$ . The histogram is totally defined by

$$\pi_k = \pi(x = k) = \frac{\sum_{i=1}^k [x_i = k]}{n}.$$

MNIST: 28x28 gray-scaled images  
each image is  $\mathbf{x} = (x_1, \dots, x_{784})$ , where  $x_i \sim \text{Be}(p_i)$ .  
 $2^{28 \times 28} - 1$  parameters to specify  $\pi(\mathbf{x})$



$$\pi(\mathbf{x}) = \pi(x_1) \cdot \pi(x_2|x_1) \cdot \dots \cdot \pi(x_m|x_{m-1}, \dots, x_1).$$

**Question:** How many parameters do we need in these cases?

$$\pi(\mathbf{x}) = \pi(x_1) \cdot \pi(x_2) \cdot \dots \cdot \pi(x_m).$$

$$\pi(\mathbf{x}) = \pi(x_1) \cdot \pi(x_2|x_1) \cdot \dots \cdot \pi(x_m|x_{m-1}).$$

## Maximum likelihood

Fix probabilistic model  $p(\mathbf{x}|\theta)$  – the set of parameterized distributions .

Instead of searching true  $\pi(\mathbf{x})$  over all probability distributions, learn function approximation  $p(\mathbf{x}|\theta) \approx \pi(\mathbf{x})$ .

### MLE problem

$$\theta^* = \arg \max_{\theta} p(\mathbf{X}|\theta) = \arg \max_{\theta} \prod_{i=1}^n p(\mathbf{x}_i|\theta) = \arg \max_{\theta} \sum_{i=1}^n \log p(\mathbf{x}_i|\theta).$$

The problem is solved with SGD.

### Requirements

- ▶ efficiently compute  $\log p(\mathbf{x}|\theta)$ ;
- ▶ efficiently compute gradient of  $\log p(\mathbf{x}|\theta)$ .

# Autoregressive model

## MLE problem

$$\theta^* = \arg \max_{\theta} p(\mathbf{X}|\theta) = \arg \max_{\theta} \prod_{i=1}^n p(\mathbf{x}_i|\theta) = \arg \max_{\theta} \sum_{i=1}^n \log p(\mathbf{x}_i|\theta).$$

## Challenge

$p(\mathbf{x}|\theta)$  could be intractable.

## Likelihood as product of conditionals

Let  $\mathbf{x} = (x_1, \dots, x_m)$ ,  $\mathbf{x}_{1:i} = (x_1, \dots, x_i)$ . Then

$$p(\mathbf{x}|\theta) = \prod_{i=1}^m p(x_i|\mathbf{x}_{1:i-1}, \theta); \quad \log p(\mathbf{x}|\theta) = \sum_{i=1}^m \log p(x_i|\mathbf{x}_{1:i-1}, \theta).$$

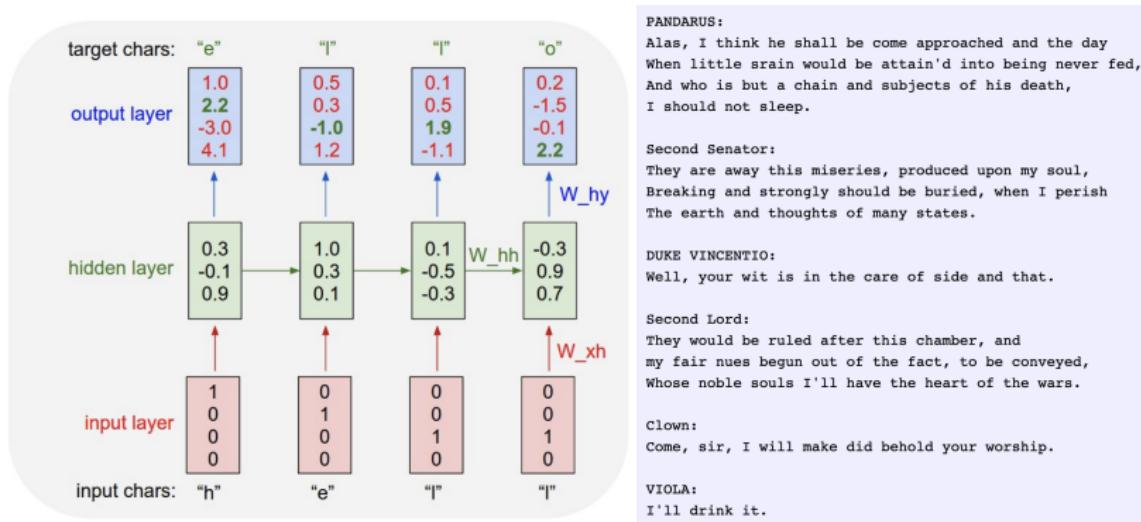
**Example:**  $p(x_1, x_2, x_3) = p(x_2) \cdot p(x_1|x_2) \cdot p(x_3|x_1, x_2).$

## Autoregressive models

$$\log p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{i=1}^m \log p(x_i|\mathbf{x}_{1:i-1}, \boldsymbol{\theta})$$

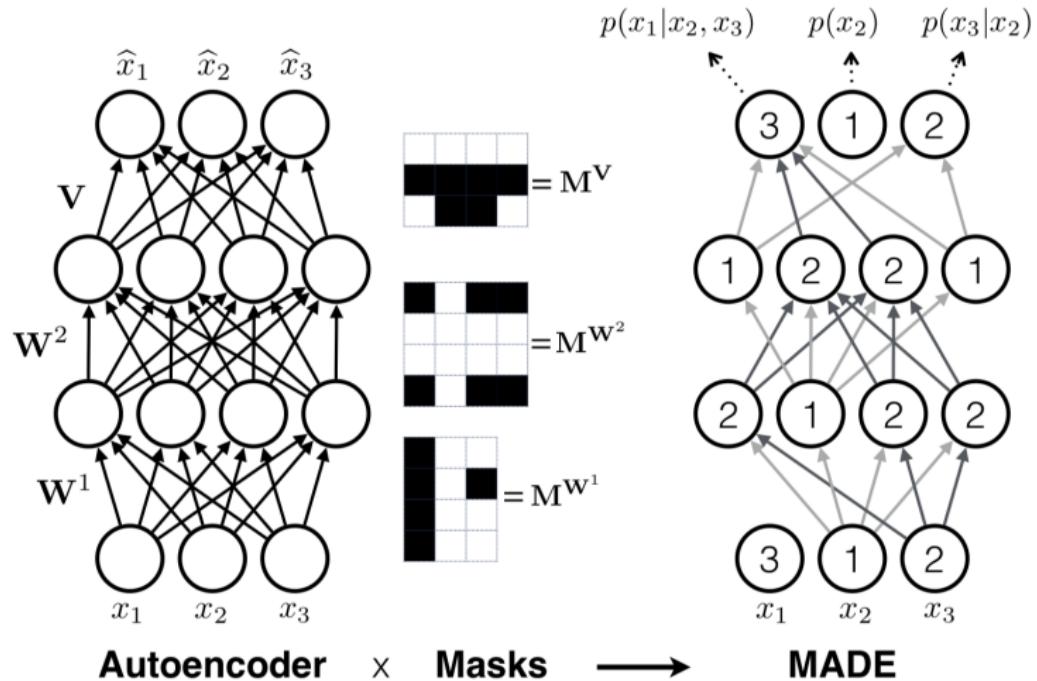
- ▶ Each conditional could be modelled by neural network.
- ▶ To extend to high dimensions share parameters across conditionals.
- ▶ Sampling is sequential.

# Char RNN



## Drawback

Sequential computation of all conditionals  $p(x_i | \mathbf{x}_{1:i-1}, \theta)$ .



# WaveNet

## Goal

Efficient generation of raw audio waveforms with natural sounds.

## Solution

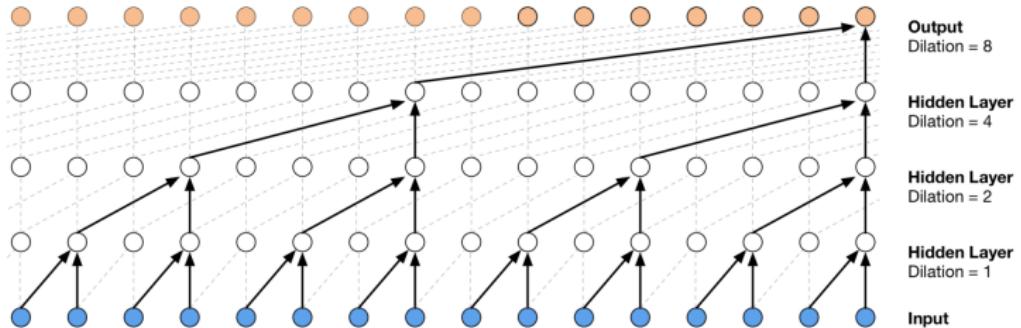
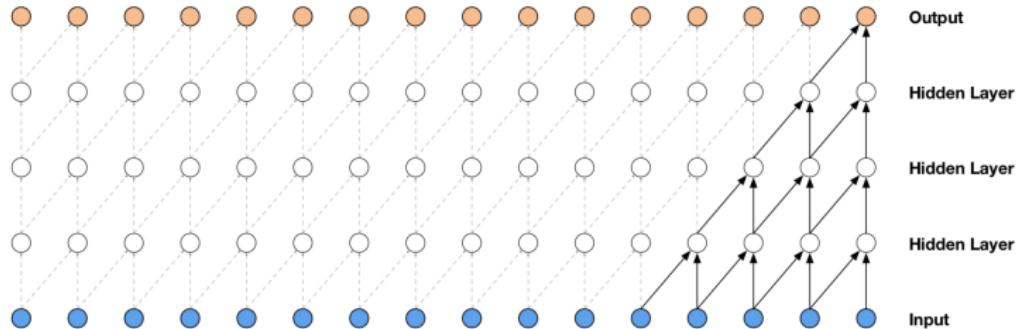
Autoregressive model

$$p(\mathbf{x}|\theta) = \prod_{t=1}^T p(x_t|\mathbf{x}_{1:t-1}, \theta).$$

The model uses causal dilated convolutions.



# WaveNet (2016)



## Goal

Model a distribution of natural images.

## Solution

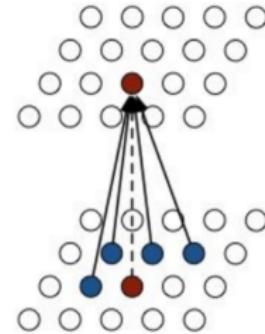
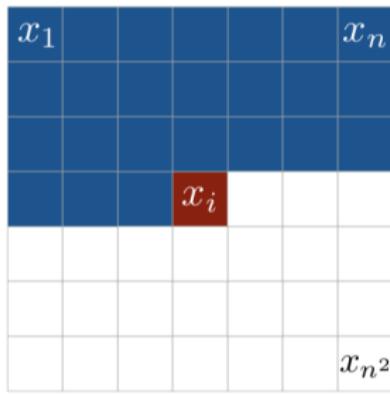
Autoregressive model

$$p(\mathbf{x}|\theta) = \prod_{i=1}^{n^2} p(x_i|\mathbf{x}_{1:i-1}, \theta).$$

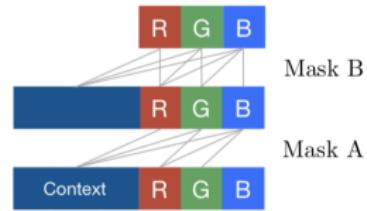
- ▶ masked convolutions;
- ▶ dependencies over RGB channels.

# PixelCNN (2016)

1	1	1
1	0	0
0	0	0



PixelCNN



## Summary

- ▶ Sampling from autoregressive models is trivial, but sequential
  - ▶ sample  $x_0 \sim p(x_0)$ ;
  - ▶ sample  $x_1 \sim p(x_1|x_0)$ ;
  - ▶ ...
- ▶ Estimating probability:

$$p(\mathbf{x}) = \prod_{i=1}^m p(x_i|\mathbf{x}_{1:i-1}).$$

- ▶ Work on both continuous and discrete data.
- ▶ There is no natural way to do unsupervised learning.

# Summary